

Possible solutions for the recap exercises

Claudius Gräbner-Radkowitsch

Table of contents

1	Packages used	1
2	CO2	1
3	Data wrangling I	3
4	Data wrangling II	4
5	Visualization and Quarto	5
6	Visualization and Quarto	6

1 Packages used

```
library(here)
library(dplyr)
library(tidyr)
library(ggplot2)
library(data.table)
```

2 CO2

We first import the raw data. Please make sure you use the **here**-package and adjust the relative paths of the following code.

Since many of the column headers were numbers (a.k.a. years), we need to make explicit that these are not values but header names. We do so by setting the optional argument `header = TRUE`:

```
co2_data_raw <- fread("co2_raw.csv", header = TRUE)
```

After inspecting the data using functions such as `str()`, `unique()` or `head()`, we first remove columns we obviously do not need and that might be irritating:

```
co2_data_tidy_1 <- co2_data_raw %>%
  select(-c(
    "Indicator Name", "Indicator Code",
    # unique() tells you there is only one indicator
    "Country Code", # Not needed
    "V69" # Sometimes such erroneous columns are part of what you download
  ))
```

We then move the year columns into rows by using `tidyr::pivot_longer()`:

```
co2_data_tidy_2 <- co2_data_tidy_1 %>%
  tidyr::pivot_longer(
    cols = -"Country Name",
    names_to = "year",
    values_to = "co2_percap")
head(co2_data_tidy_2)
```

```
# A tibble: 6 x 3
  `Country Name` year  co2_percap
  <chr>          <chr>    <dbl>
1 Aruba         1960      NA
2 Aruba         1961      NA
3 Aruba         1962      NA
4 Aruba         1963      NA
5 Aruba         1964      NA
6 Aruba         1965      NA
```

We see that the year column is still a **character**. So we transform it into a **double** to then filter the years. We can also filter for the required countries within the same function call and then rename the column:

```
co2_data_tidy_3 <- co2_data_tidy_2 %>%
  mutate(year = as.double(year)) %>%
  filter(
    year >= 2000, year<=2020,
    `Country Name` %in% c(
      "South Africa", "United States", "Sub-Saharan Africa",
      "European Union", "Germany", "China")
  ) %>%
  rename(country = `Country Name`)
```

We could have done everything in one call as well:

```
co2_data_tidy <- co2_data_raw %>%
  select(-c(
    "Indicator Name", "Indicator Code",
    # unique() tells you there is only one indicator
    "Country Code", # Not needed
    "V69" # Sometimes such erroneous columns are part of what you download
  )) %>%
  tidyr::pivot_longer(
    cols = -"Country Name",
    names_to = "year",
    values_to = "co2_percap") %>%
  mutate(year = as.double(year)) %>%
  filter(
    year >= 2000, year<=2020,
    `Country Name` %in% c(
      "South Africa", "United States", "Sub-Saharan Africa",
      "European Union", "Germany", "China")
  ) %>%
  rename(country = `Country Name`)
```

Then think about a useful location to store the data and do something like:

```
fwrite(co2_data_tidy, file = here("data/tidy/co2_tidy.csv"))
```

3 Data wrangling I

Please make sure you use the `here`-package and adjust the relative paths of the following code:

Compute, for each country, the percentage change of the spending from the year 2010 to the year 2020 and save this as a variable called `perc_change`.

```
educ_exercise_data_raw <- fread("education_income.csv")

educ_exercise_data <- educ_exercise_data_raw %>%
  dplyr::select(-c("income", "GDPpc")) %>%
  dplyr::filter(year %in% c(2010, 2020)) %>%
  tidyr::pivot_wider(
    names_from = "year",
    values_from = "EducationSpending"
  ) %>%
  dplyr::mutate(
    perc_change = ((`2020` - `2010`) / `2010`) * 100
  ) %>%
  dplyr::filter(!is.na(perc_change))
head(educ_exercise_data)
```

```
# A tibble: 6 x 4
  iso3c `2010` `2020` perc_change
  <chr>   <dbl>   <dbl>         <dbl>
1 ALB     3.41     3.34          -2.07
2 AND     2.98     2.86          -4.05
3 AGO     3.42     2.74         -19.8
4 ARG     5.02     5.28           5.18
5 ARM     3.25     2.71         -16.7
6 AUS     5.54     5.61           1.27
```

Then think about a useful location to store the data and do something like:

```
data.table::fwrite(
  x = educ_exercise_data,
  file = here("data/tidy/educ_perc_change.csv"))
```

4 Data wrangling II

We use `educ_exercise_data_raw` as imported above as a starting point and proceed as follows:

Compute for each income group the average expense of education over the whole period. Make sure missing values are ignored.

Save the new data set under a useful name in an adequate location.

```
educ_exercise_summarized <- educ_exercise_data_raw %>%  
  dplyr::summarise(  
    EducExpense_avg = mean(EducationSpending, na.rm = TRUE),  
    .by = "income")  
educ_exercise_summarized
```

	income	EducExpense_avg
1	Low income	3.342121
2	Upper middle income	4.723520
3	Lower middle income	4.534556
4	High income	4.644050

Then think about a useful location to store the data and do something like:

```
data.table::fwrite(  
  x = educ_exercise_summarized,  
  file = here("data/tidy/educ_perc_income-groups.csv"))
```

5 Visualization and Quarto

The quarto header should look like this:

```
title: "Sessions 12 and 13: Recap and Practice"  
author: "Claudius Gräbner-Radkowitzsch"  
format:  
  html:  
    number-sections: true  
    table-of-contents: true  
    toc-location: body  
execute:  
  echo: false  
  warning: false  
  message: false
```

6 Visualization and Quarto

To read in the data set do something as the following, but make sure you are using the `here`-package and set the path accordingly.

```
child_mortality <- data.table::fread("child_mortality.csv")
head(child_mortality)
```

	iso3c	year	ChildMortality	GDPpc
	<char>	<int>	<num>	<num>
1:	AFG	2017	64.6	2096.093
2:	AFG	2014	73.4	2110.830
3:	AFG	2016	67.2	2023.835
4:	AFG	2012	80.3	1958.448
5:	AFG	2021	55.7	1673.144
6:	AFG	2007	100.0	1287.064

To summarize the data:

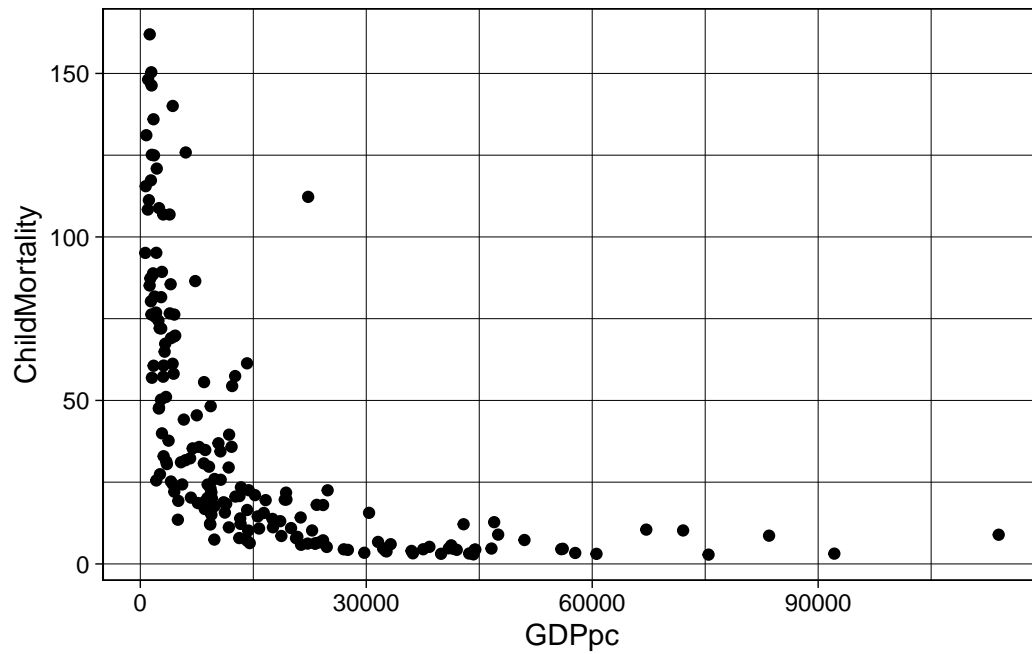
```
child_mortality_summarized <- child_mortality %>%
  dplyr::summarise(
    ChildMortality = mean(ChildMortality, na.rm = TRUE),
    GDPpc = mean(GDPpc, na.rm = TRUE),
    .by = "iso3c")
head(child_mortality_summarized)
```

	iso3c	ChildMortality	GDPpc
1	AFG	88.459091	1660.568
2	ALB	15.031818	9437.101
3	DZA	29.486364	11735.174
4	ASM	NaN	NaN
5	AND	4.718182	NaN
6	AGO	125.831818	6029.127

We can then directly create a simple scatter plot:

```
ggplot2::ggplot(
  data = child_mortality_summarized,
  mapping = aes(x = GDPpc, y = ChildMortality)
) +
```

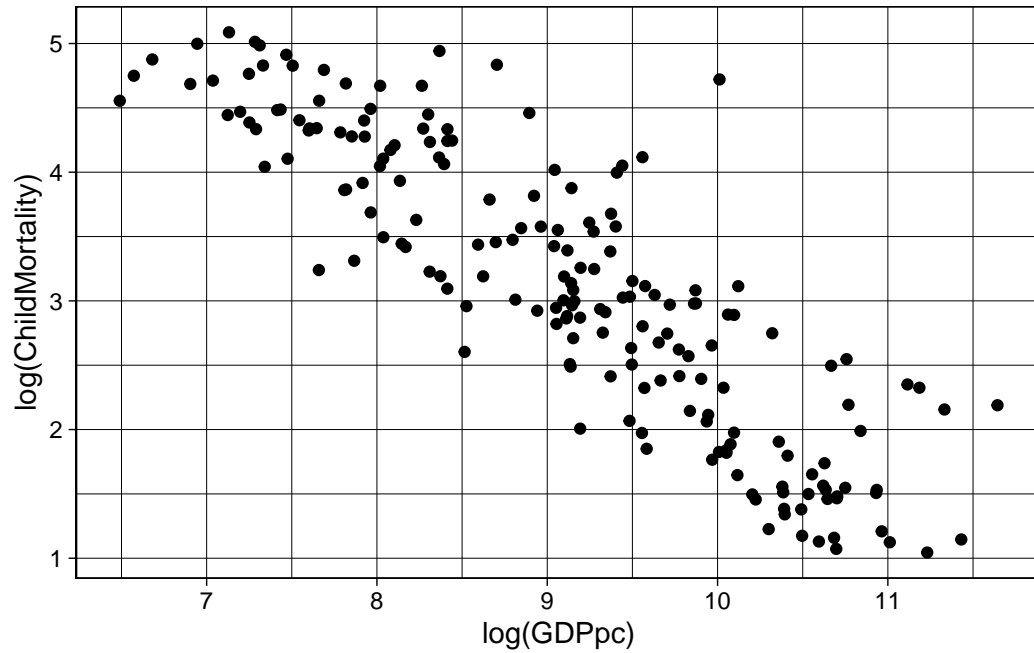
```
geom_point() +  
theme_linedraw()
```



We see a clear non-linear relationship.

We now plot the data in logarithms. You can do this by changing the underlying data, rescale an axis, or make the change directly in the `data` argument of `ggplot2::ggplot()`:

```
ggplot2::ggplot(  
  data = child_mortality_summarized,  
  mapping = aes(x = log(GDPpc), y = log(ChildMortality))  
) +  
  geom_point() +  
  theme_linedraw()
```



The relationship now becomes almost linear. This is typical for relationships that are exponential. We can say: an increase in GDP per capita by one percent is on average associated with a reduction of child mortality by 0.83 per cent (the latter value is given by a regression, but we come to this later).